

Big data service systems: Models and Platforms

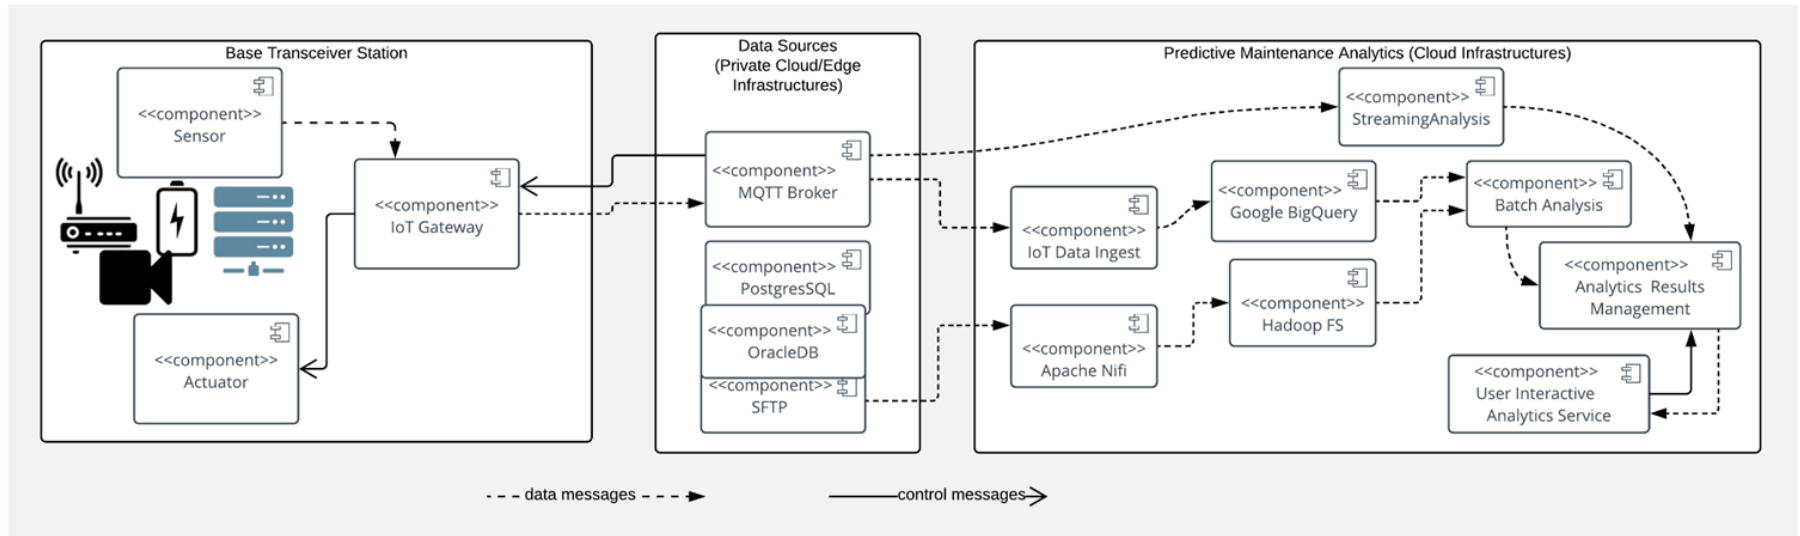
Hong-Linh Truong
Faculty of Informatics, TU Wien

hong-linh.truong@tuwien.ac.at
www.infosys.tuwien.ac.at/staff/truong@linhsolar

- Data analytics within a single system
- Data analytics across multiple systems
- APIs management and big data systems
- Principles of elasticity for advanced service-based data analytics

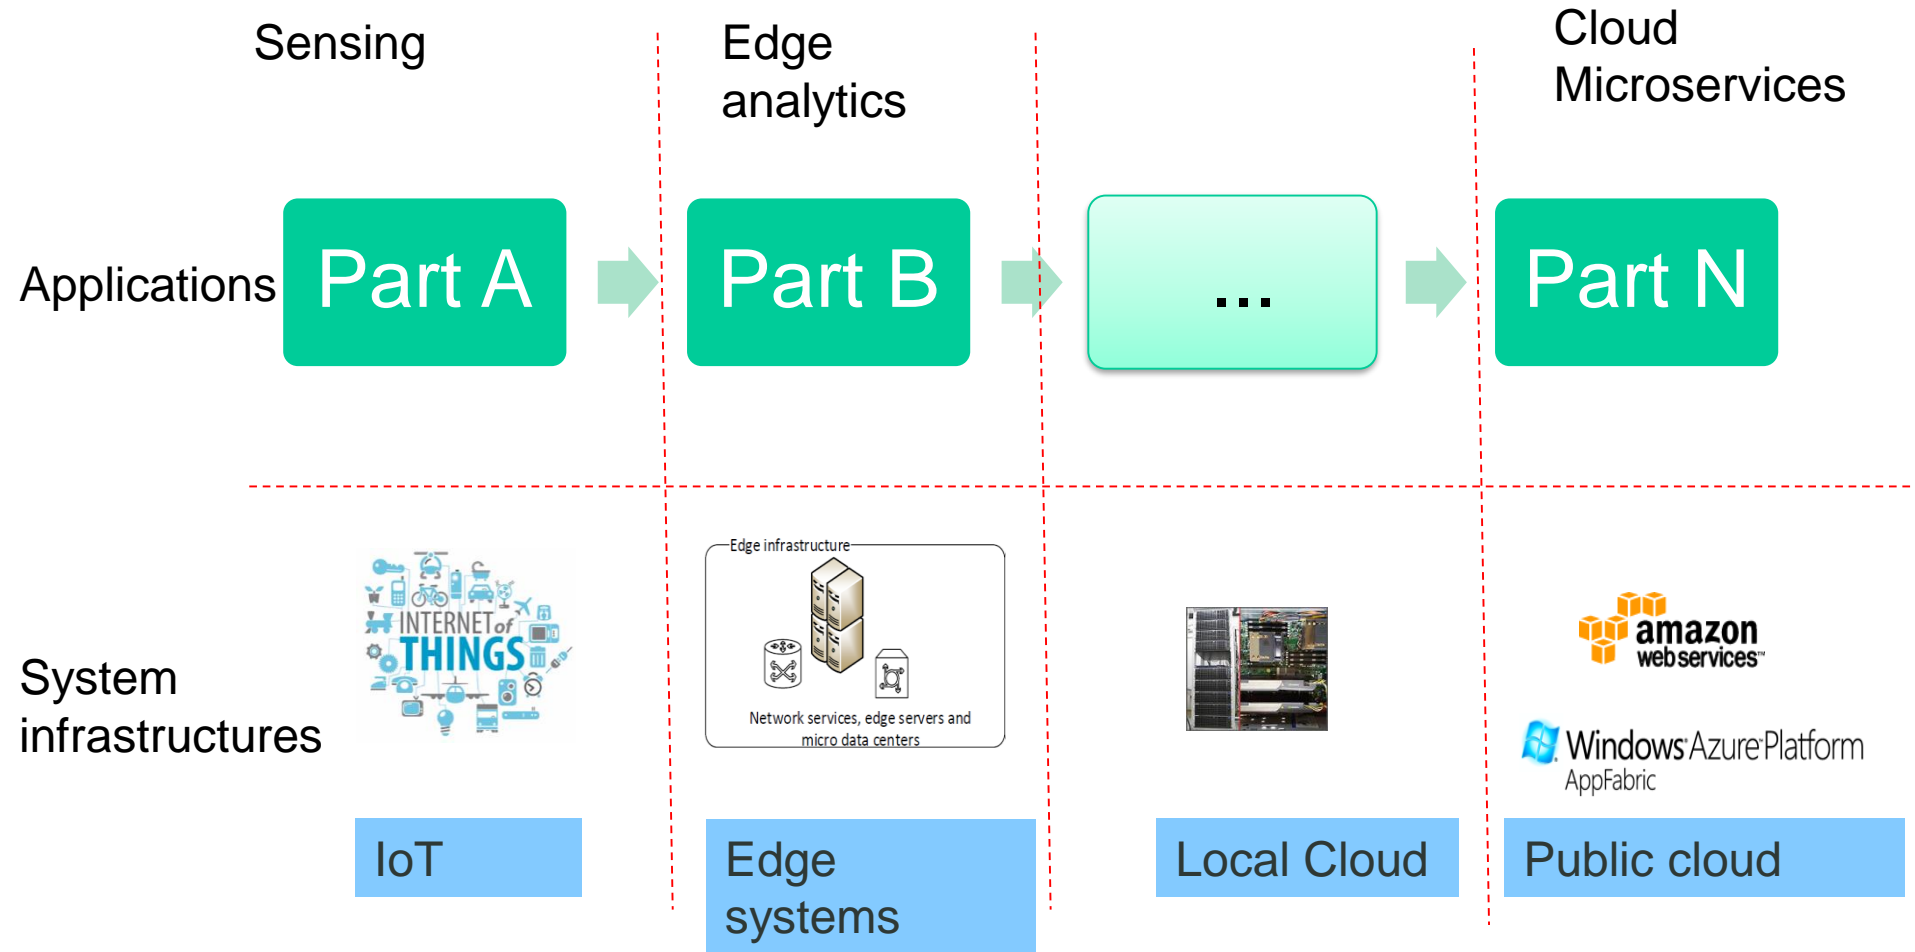
Advanced service-based analytics – which are fundamental engineering questions?

Scenario: Predictive Maintenance in Telcos



- Complex types of data
- Various services
- Complex analytics/data processing algorithms

Advanced service-based data analytics -- fundamental concepts




Design questions

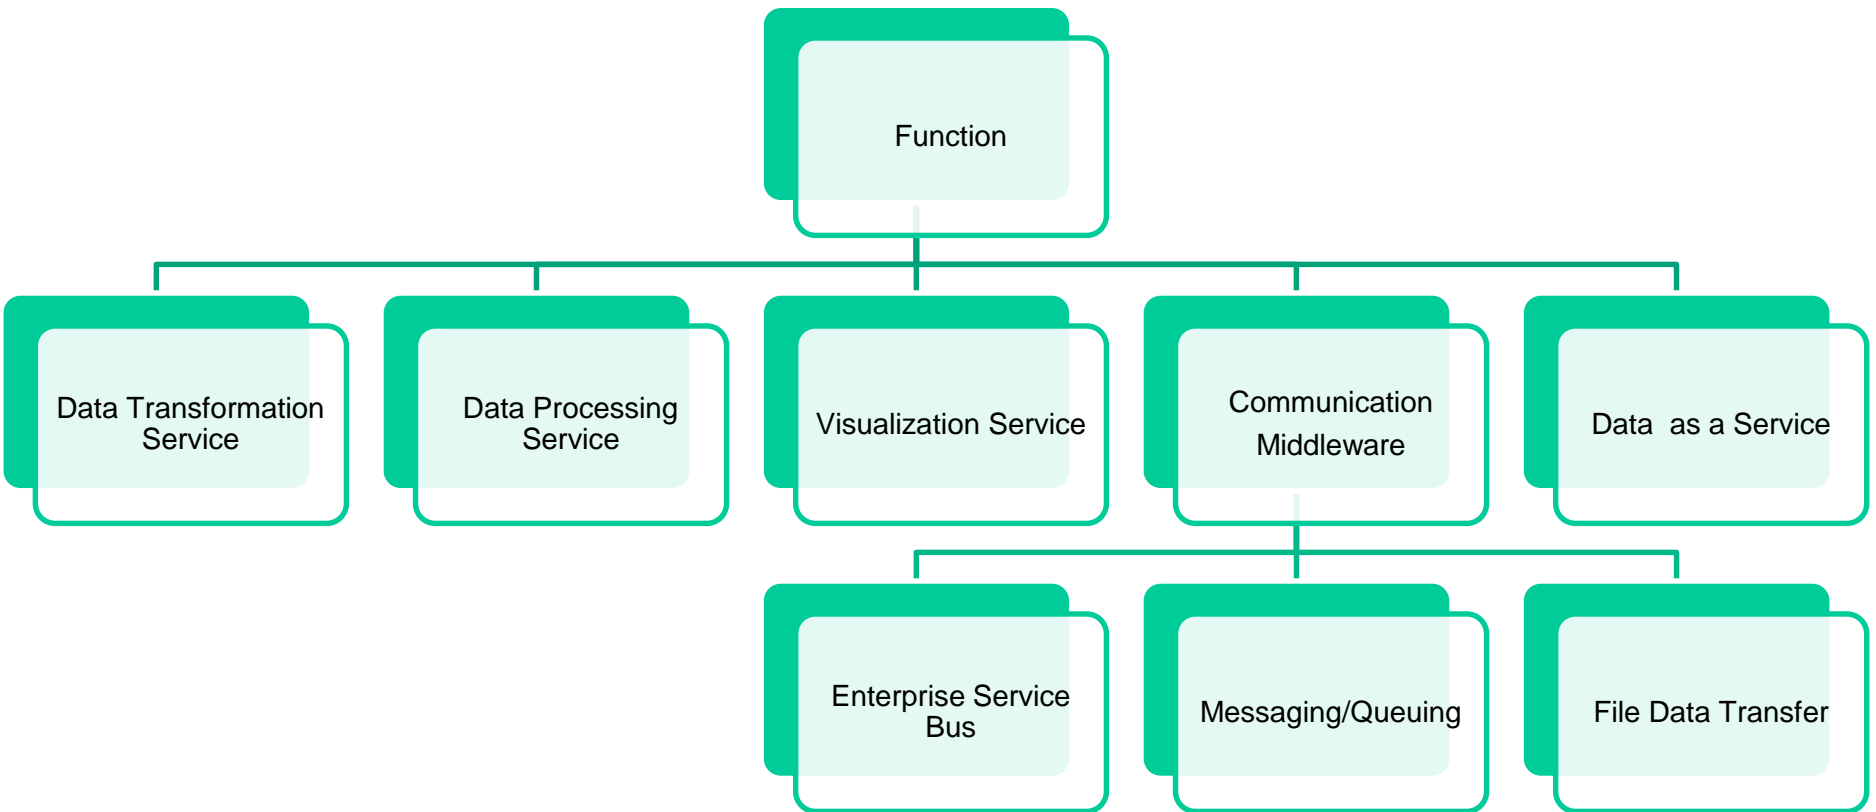
Part = a (composite) (micro)services/components

- Which system **infrastructures** are used?
- Which **interfaces/APIs** are suitable for services?
- Which **programming models** are used within services?
- Which **non-functional parameters** are important and how to measure them?

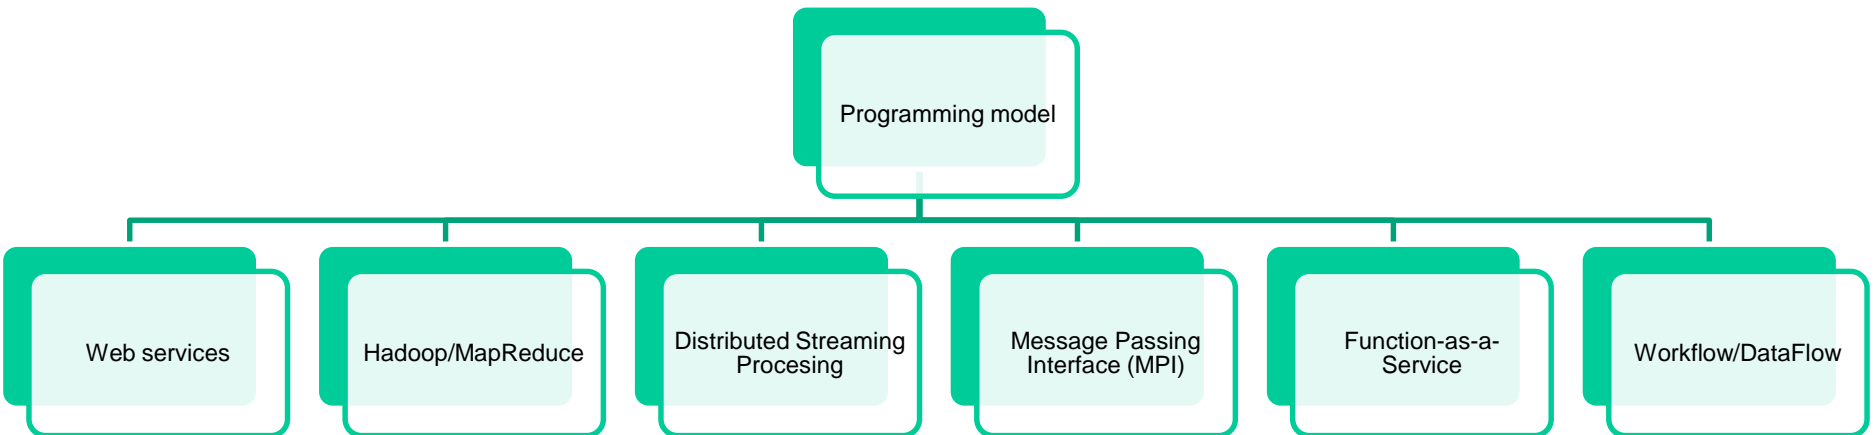
Fundamental concepts – system infrastructure unit



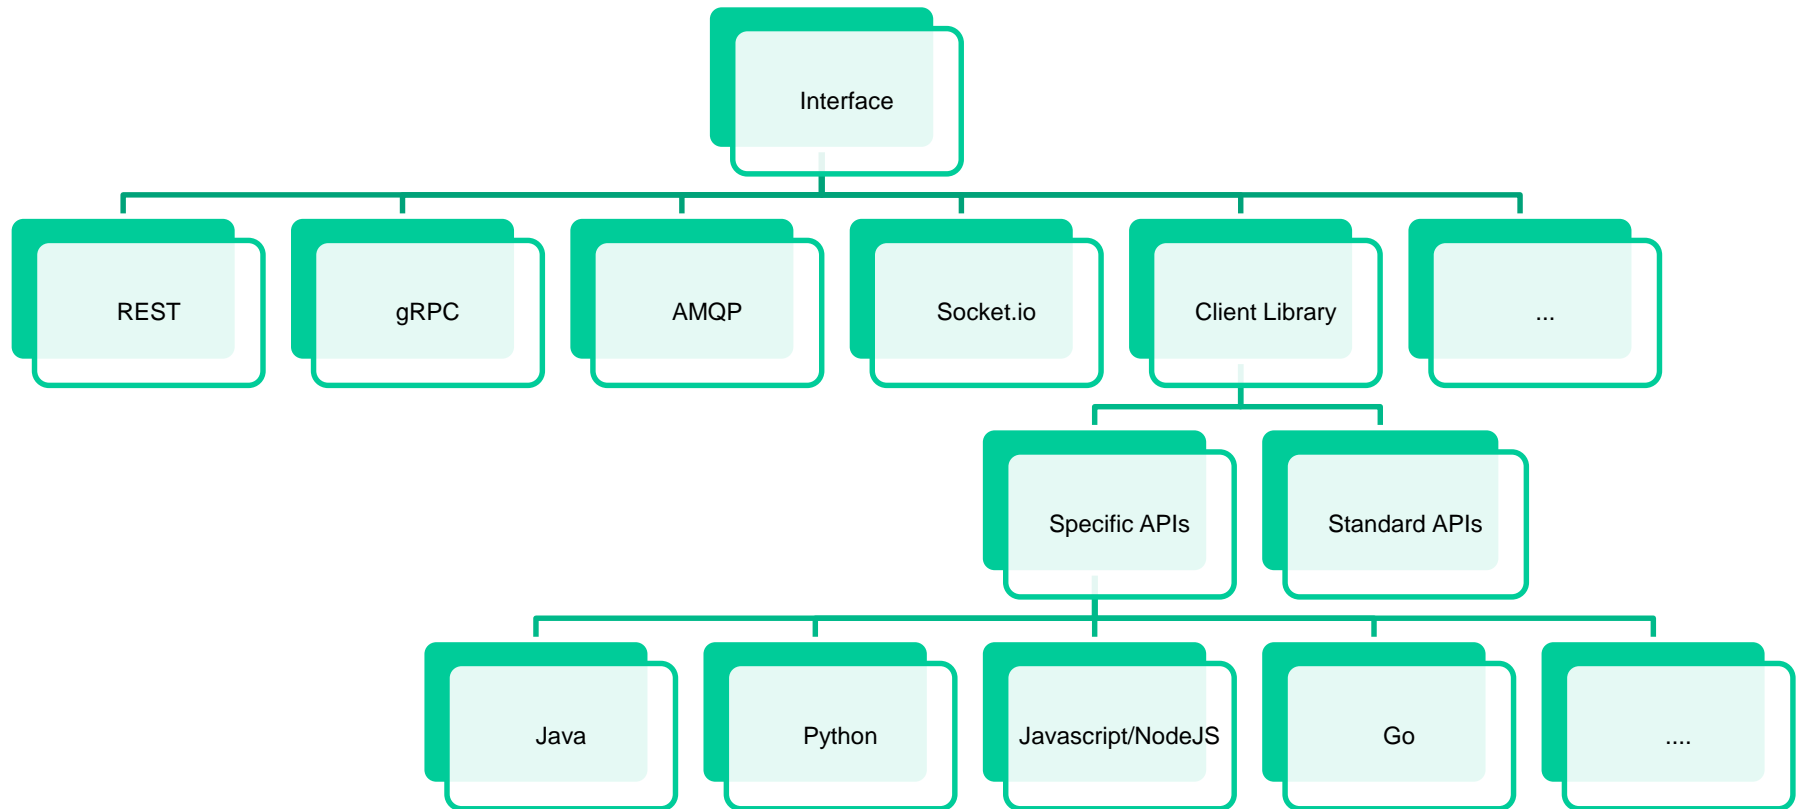
Fundamental concepts – unit functions



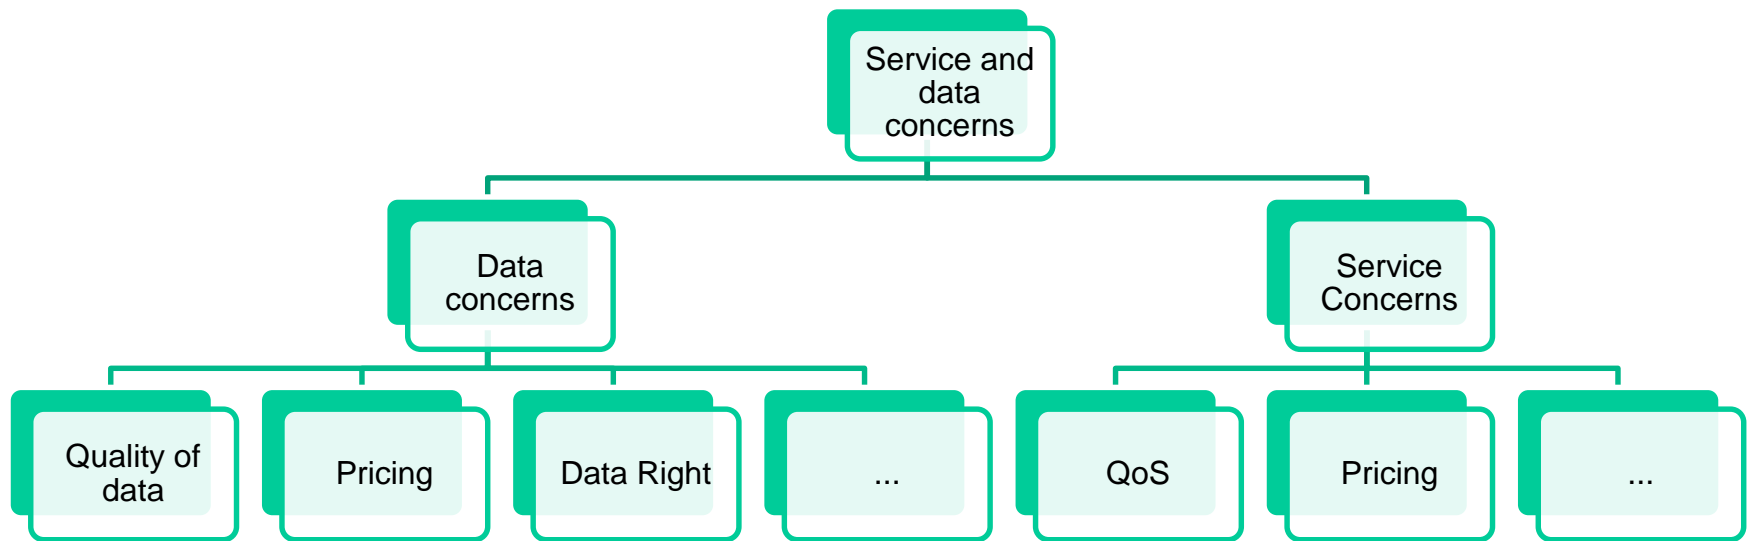
Fundamental concepts – programming model within units



Fundamental concepts – interfaces between services



Fundamental concepts – services and data concerns



You see we need to deal with many techniques and frameworks


WE NEED TO START FROM DATA ANALYTICS WITHIN A SINGLE SYSTEM

What is our understanding about a single system?

Location and enterprise boundary?

Within a virtual infrastructure owned by a single organization?

Data analytics within a single (technical) system



- In a single domain
 - Tightly coupled computing infrastructures
 - E.g., in the same cloud
 - Computation and data are close
 - Several concerns can be bypassed
 - „Single system“ does not mean „simple“: they can be complex

Data analytics within a single system – some examples

Message Passing
Interface (MPI) + Cluster-
based File system

Parallel Database
(SQL/NonSQL)

Big Query

Azure HDInsight

Hadoop + HDFS

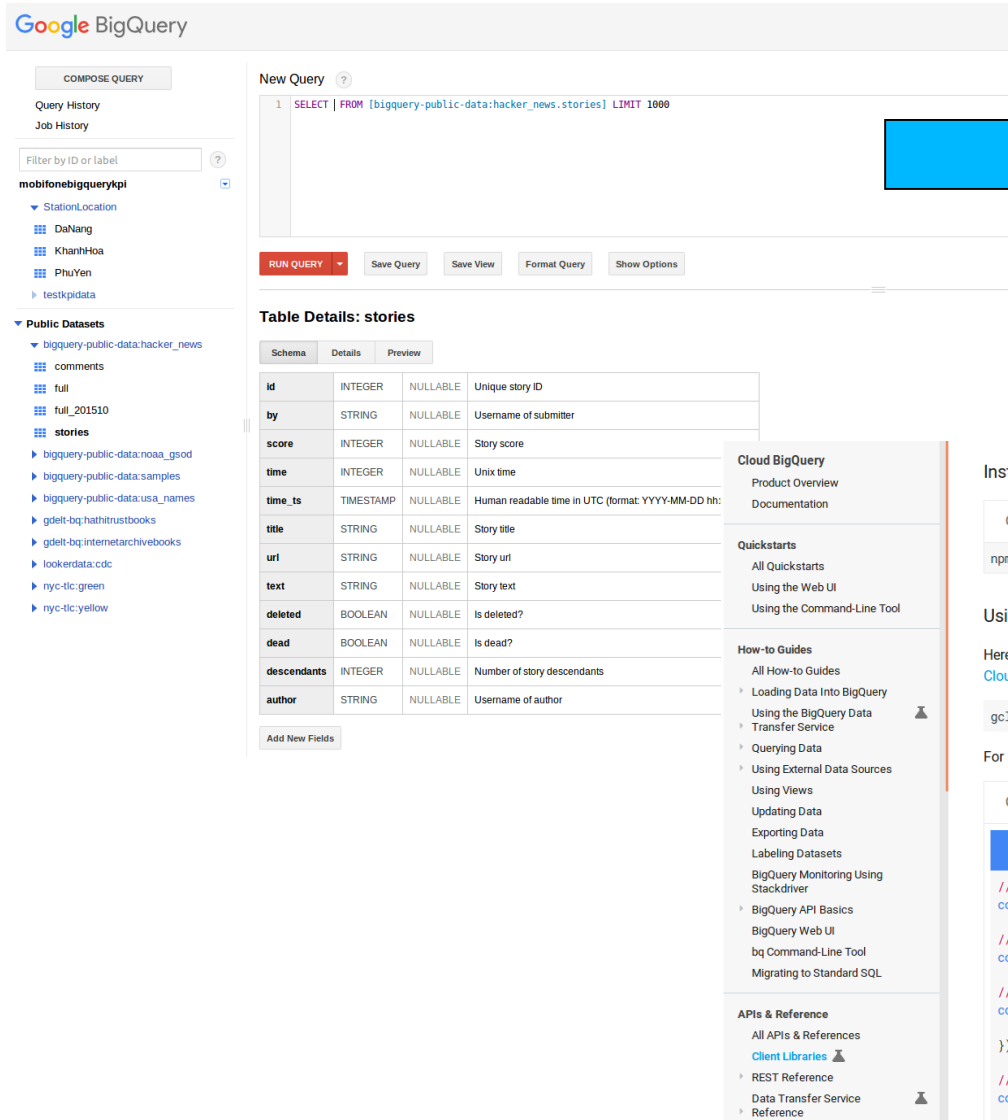
Apache Spark

Amazon RedShift

Scientific/Business
Workflow

A short, good overview in Chapter 6: Cloud Programming and Software Environments, Book: Distributed and Cloud Computing – from Parallel Processing to the Internet of Things, Kai Hwang, Geoffrey C. Fox and Jack J Dongarra, Morgan Kaufmann, 2012


Example - BigQuery (1)



The screenshot shows the Google BigQuery web interface. On the left, there's a sidebar with navigation options like 'COMPOSE QUERY', 'Query History', and 'Job History'. Below that, a list of datasets is shown, including 'mobifonebigquerykpi' and 'Public Datasets'. The main area displays a 'New Query' with a SQL query: `SELECT * FROM [bigquery-public-data:hacker_news.stories] LIMIT 1000`. Below the query editor, there are buttons for 'RUN QUERY', 'Save Query', 'Save View', 'Format Query', and 'Show Options'. To the right of the query editor, the 'Table Details: stories' are shown in a tabular format with columns like 'id', 'by', 'score', 'time', 'time_ts', 'title', 'uri', 'text', 'deleted', 'dead', 'descendants', and 'author'. A sidebar on the right contains links for 'Cloud BigQuery', 'Quickstarts', 'How-to Guides', and 'APIs & Reference'.



Google
BigQuery



Installing the client library

```
C#  GO  JAVA  NODE.JS  PHP  PYTHON  RUBY

npm install --save @google-cloud/bigquery
```

Using the client library

Here's an example of how to use the client library. To run it on your local workstation you must first install the [Google Cloud SDK](#) and authenticate by running [the following command](#):

```
gcloud auth application-default login
```

For information about authenticating in other environments, see the [Google Cloud Platform Auth Guide](#).

```
C#  GO  JAVA  NODE.JS  PHP  PYTHON  RUBY

VIEW ON GITHUB  FEEDBACK

// Imports the Google Cloud client library
const BigQuery = require('@google-cloud/bigquery');

// Your Google Cloud Platform project ID
const projectId = 'YOUR_PROJECT_ID';

// Instantiates a client
const bigquery = BigQuery({
  projectId: projectId
});

// The name for the new dataset
const datasetName = 'my_new_dataset';
```

Example – BigQuery: complexity





Figure 1: BigQuery structural overview

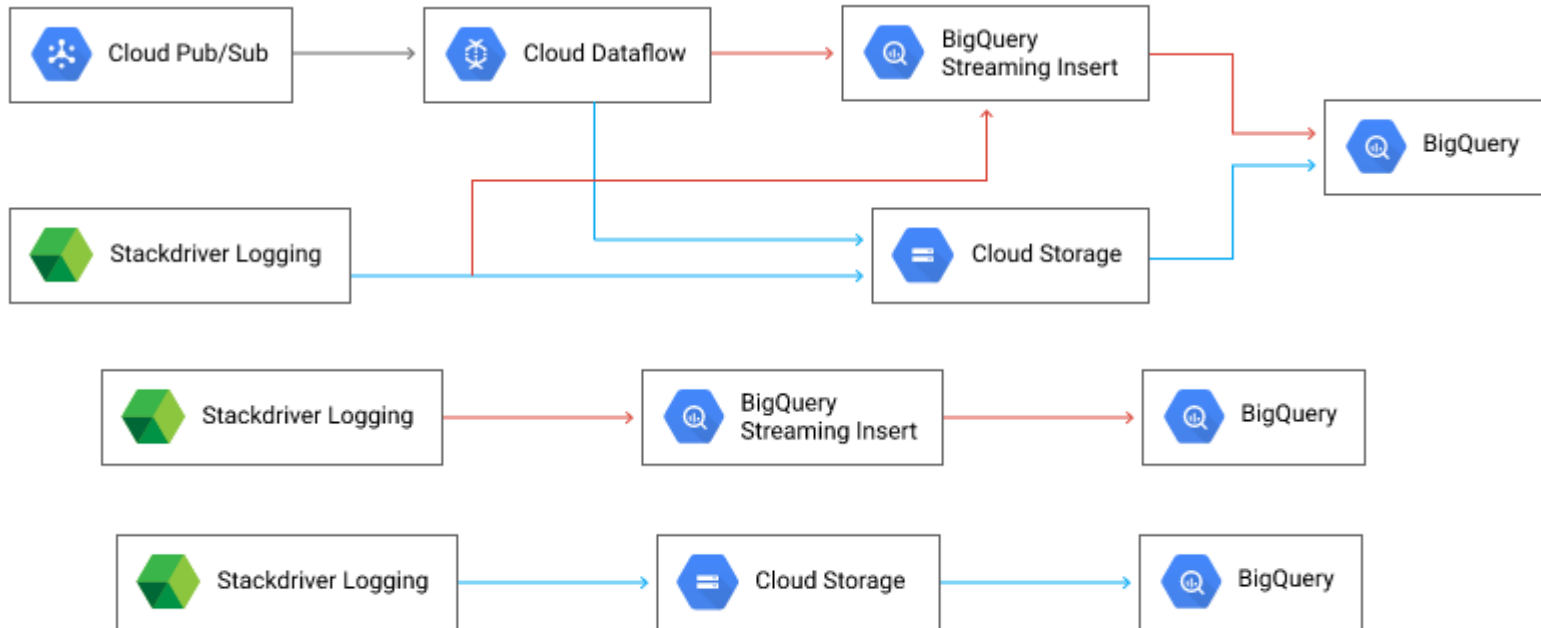
Source <https://cloud.google.com/solutions/bigquery-data-warehouse>

Example – BigQuery: complexity



Source: <https://cloud.google.com/blog/big-data/2016/01/bigquery-under-the-hood>

Example – BigQuery: integration patterns



Source: <https://cloud.google.com/solutions/architecture/optimized-large-scale-analytics-ingestion>


Example - Hadoop

```

truong@bachphu-spark-m: ~
The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
truong@bachphu-spark-m:~$ ls
aa  linh.csv  spark-warehouse  tt.py
truong@bachphu-spark-m:~$ hadoop fs -ls /user/truong
17/05/18 21:03:20 INFO gcs.GoogleHadoopFileSystemBase: GHFS version: 1.6.0-hadoop2
Found 4 items
drwxr-xr-x  - truong hadoop      0 2017-05-17 14:29 /user/truong/.sparkStaging
-rw-r--r--  2 truong hadoop    8945 2017-05-12 07:42 /user/truong/aa
drwxr-xr-x  - truong hadoop      0 2017-05-12 07:40 /user/truong/output
-rw-r--r--  2 truong hadoop    8945 2017-05-12 07:41 /user/truong/part-r-00000-9f88111c-f139-40e5-ac06-53a6e283cd40.csv
truong@bachphu-spark-m:~$ hadoop fs -copyFromLocal aa /user/truong/test.csv
17/05/18 21:04:00 INFO gcs.GoogleHadoopFileSystemBase: GHFS version: 1.6.0-hadoop2
truong@bachphu-spark-m:~$


```



Hadoop File Systems

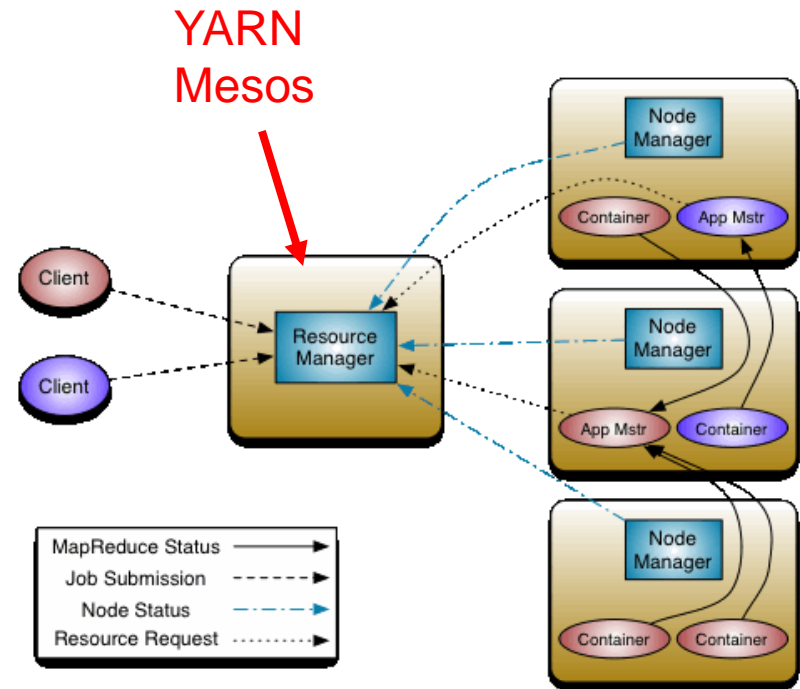
Example – Hadoop: complexity

- Distributing data into multiple nodes/machines is the key! Why?
- Hadoop provides a parallel file system – Hadoop File Systems
 - Deal with hardware failures, support data locality, streaming data access
 - Like traditional file systems with new features for big data
- Key principles:



Example – Hadoop: complexity

- Several computers are used to setup Resource Manager and Node Manager
- You write the tasks and you submit the tasks



Source: <http://hadoop.apache.org/docs/current/hadoop-yarn/hadoop-yarn-site/YARN.html>

Example – Hadoop: simple

```
spark =
SparkSession.builder.appName("sp_AlarmTypePerday").getOrCreate()
df
=spark.read.csv("hdfs://test/Alarm_nodeB_DN_9_Jan.csv",header=True,
inferSchema=True)
newdf = df.select(['Alarm Number','Started','Canceled'])
newdf.show()
```

Submission
command
line
(your local
machine)

```
gcloud dataproc jobs submit pyspark
--cluster cluster-spark Test.py
```

Google
cloud

cluster-spark
(Google spark installation with n
nodes)

```
spark-submit --master
spark://master-node:7077 Test.py
```

master-node

Slave
node

Slave
node

Slave
node

Slave
node

Slave
node

Similar examples

- With ElasticSearch, MongoDB, Cassandra, etc. within a single system → they are very large and scalable!

But when are they not enough? When are they not suitable for us?

Bigdata-as-a-Service for ASE

- How do we use them?
 - Develop services for bigdata-as-a-service
 - Develop services utilizing bigdata-as-a-service
- Important issues
 - Design concepts and architectural patterns
 - Microservices
 - Data access protocols and performance patterns
- Abstraction
 - Data service units and data concerns
 - Performance concerns


Data Analytics Unit: Characteristics

A blue oval containing the text 'Data Analytics Unit' in white.

Data
Analytics Unit

- Can be simple or complex
 - E.g., a python program based on scikit-learn or a pySpark program or an Airflow workflow
- Can be written in different program languages
- Can be deployed and run “as a service”
 - Clear input & output

Data analytics across multiple systems – data service units



Interface

- Read/write data via direct , low-level read/write via IO

System

- Cluster or cluster of clusters
- Can be very large

Programming model

- Usually parallel processing


NFS

Lustre

Hadoop File System

Google file system

Data analytics across multiple systems – data service units



Interface

- Direct data transfer via REST APIs

System

- Decouple between analytics and storage


Programming model

- May require middleware for data transfer
- Request via REST
- Real data transfer done by external middleware
- A rich set of programming models can be used

Amazon S3
(SOAP/REST API)

Google Storage Service
(REST API)

Data analytics across multiple systems – data service units



Interface

- REST APIs
- Mainly for commands and results

System

- Decouple between analytics unit and database
- Database as a service can be very large

Programming model


- Analytics can be done at both sides
- Analytic units can use any programming models
- Database-as-a-service can perform a lot of analytics
- Parallel database operations

Technology

MongoDB/MongoLab
Amazon DynamoDB
Amazon SimpleDB
Cloudant Data

SkySQL
Amazon RDS
Microsoft SQL Azure
Clustrix DBaaS

Data analytics across multiple systems – data service units



Interface

- Data transfer can be uni or bi-direction
- Streaming data protocols

System

- Both systems for DaaS and for analytics units can be very large

Programming model

- Can be any

Technology

StormMQ, RabbitMQ,
CloudMQTT, Google Data
Hub, Azure Data Hub, ...


WHY SHOULD ANALYTICS UNITS BE „CLOSED“ TO DATA UNITS?

WHICH CONCERNS COULD BE IGNORED IN SINGLE SYSTEM DATA ANALYTICS?

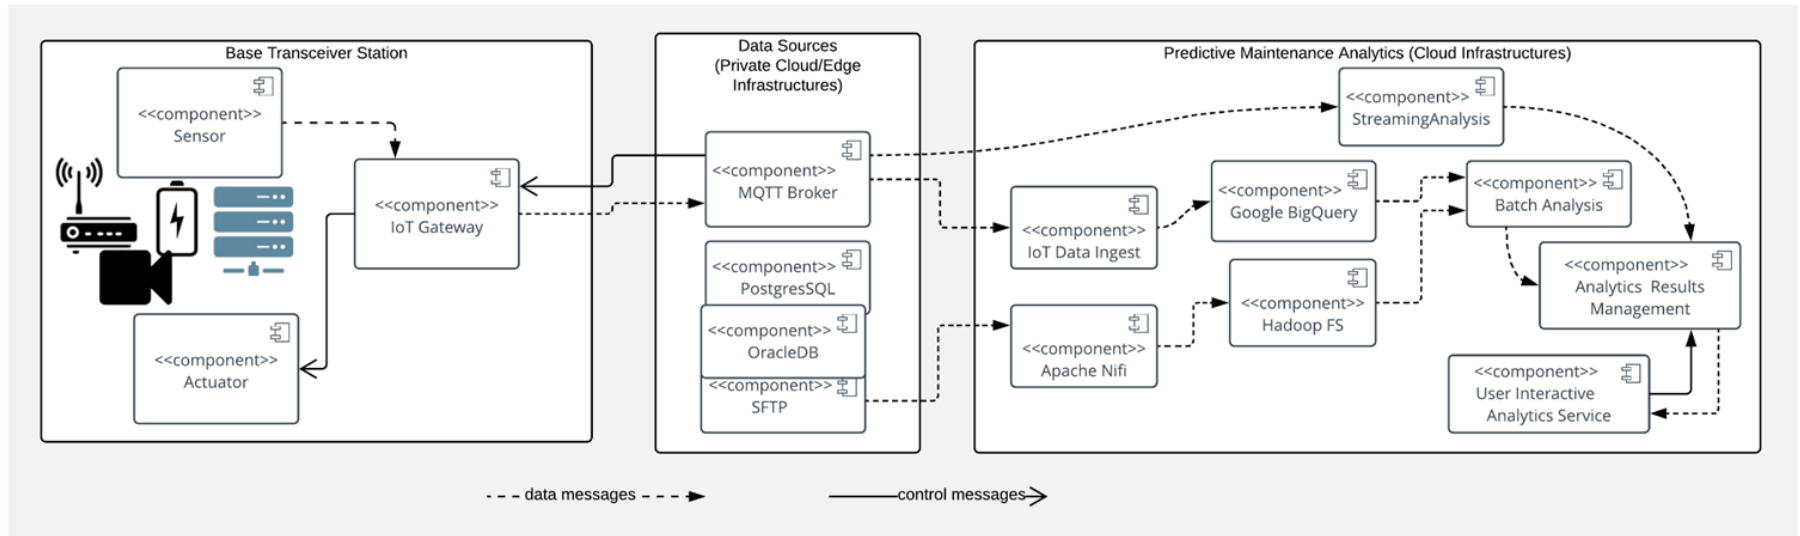
**WHICH ARE THE ISSUES THAT WE NEED
TO CONSIDER WHEN OUR DATA UNITS
ARE IN DIFFERENT SYSTEMS?**

Data analytics across multiple systems – design choice

- Programming models for data analytics service
- Data service units
- Supporting middleware units



Data analytics across multiple systems - example

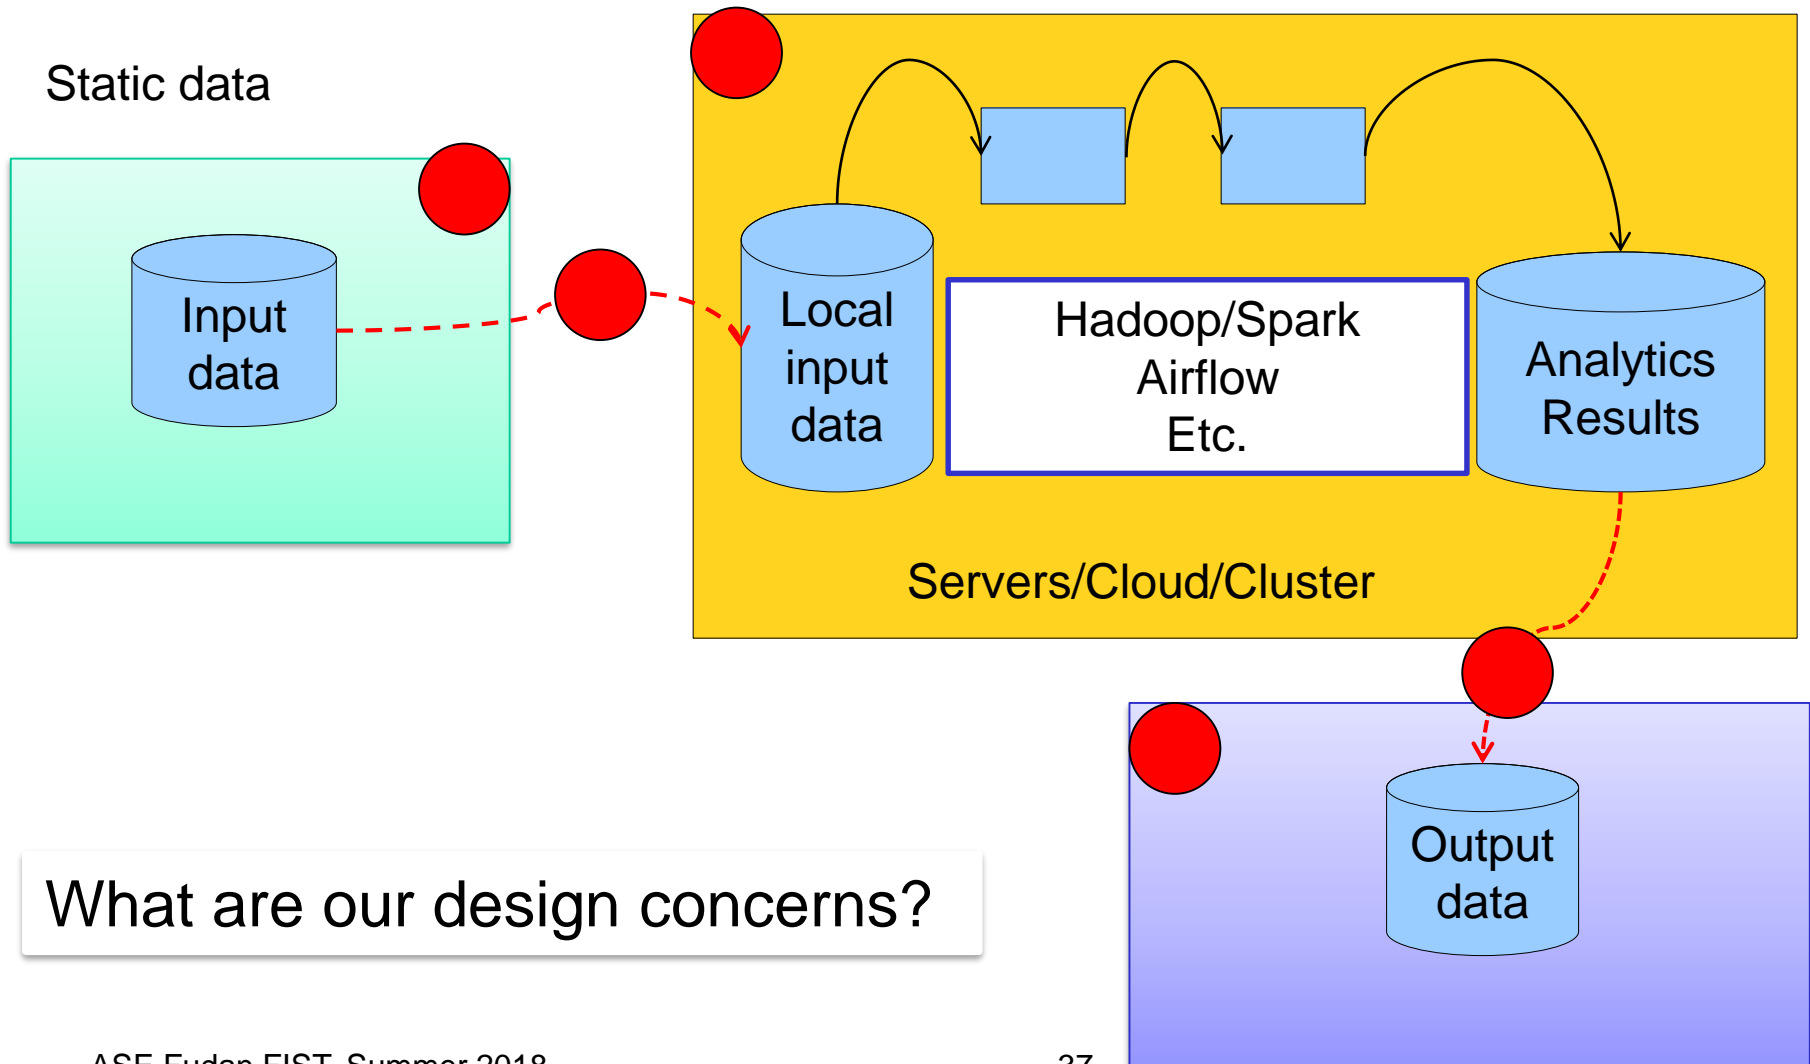


How many systems?

Programming languages?

Type of data?

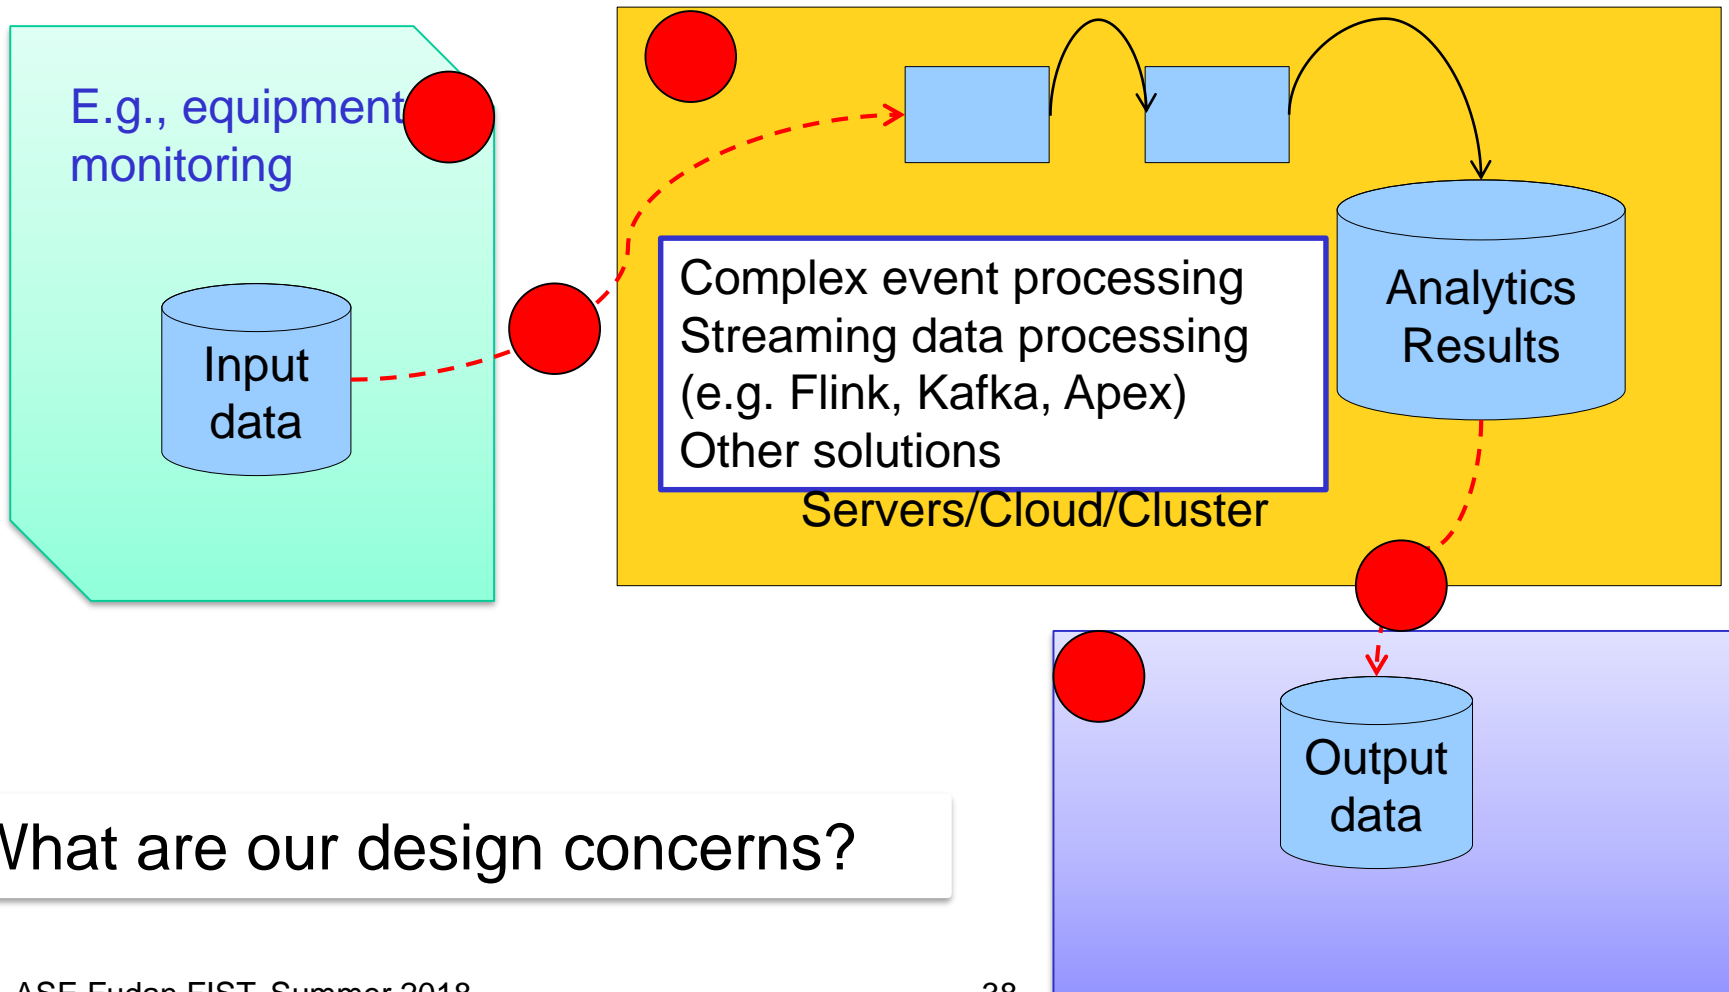
Data analytics across multiple systems – programming models (1)



What are our design concerns?

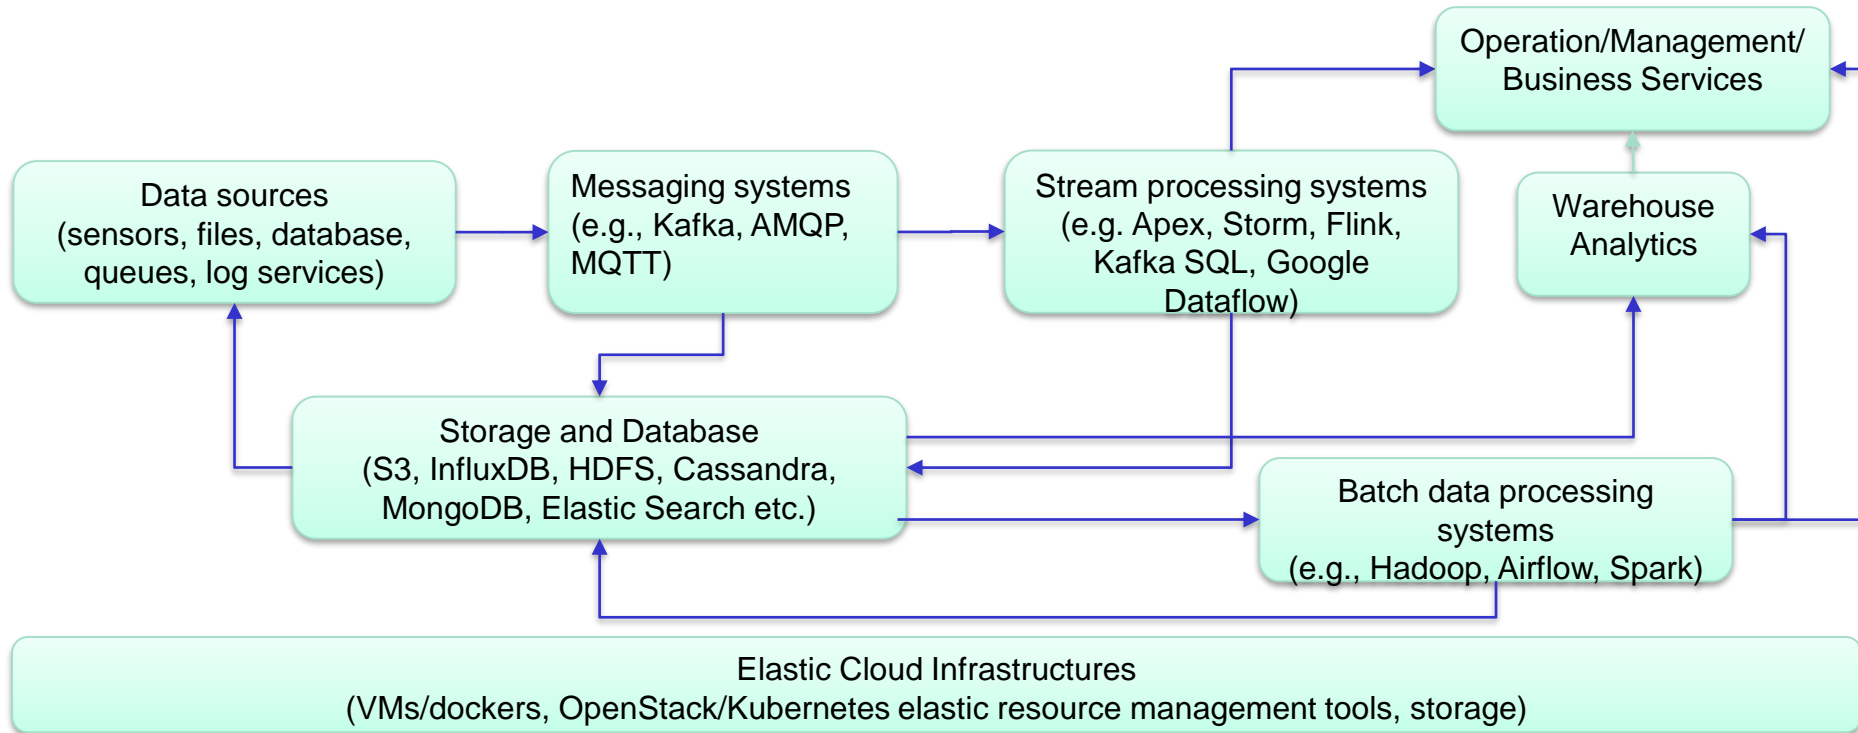
Data analytics across multiple systems – programming models (2)

Near-realtime data



What are our design concerns?

Cloud services and big data analytics



Very complex problems due to software complexity, infrastructures management and service providers

Case studies

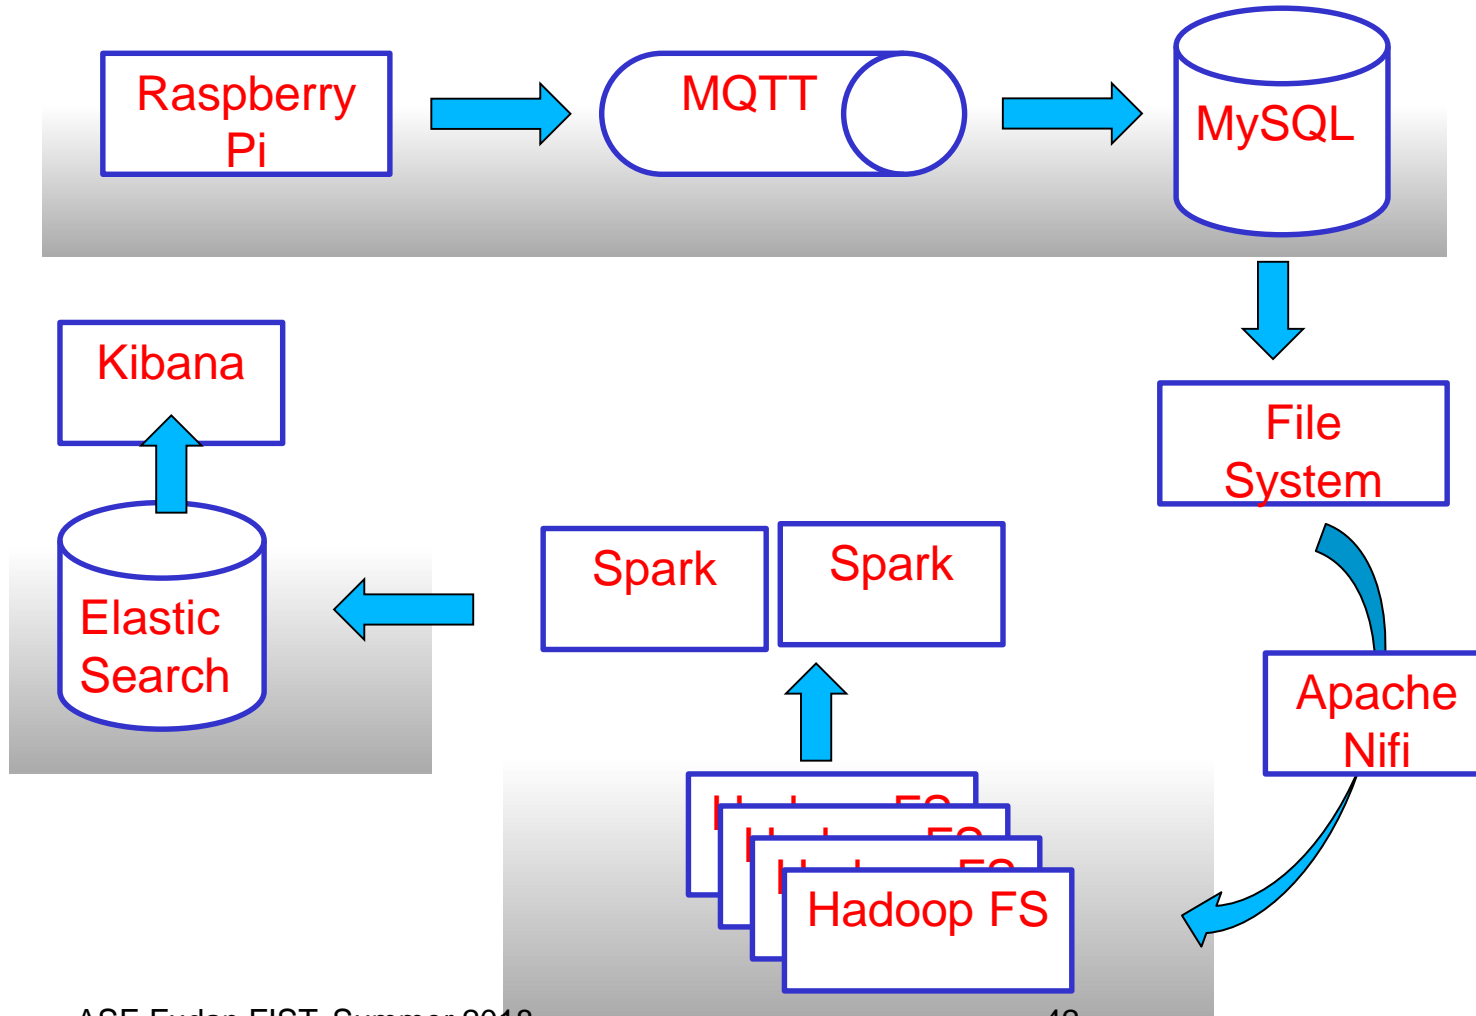
- Monitoring equipment and environments
 - Electricity, temperature, air conditioner breakdown, etc.
- Using MQTT and MySQL



- **Requirements:**
 - Now would like to do big data analytics (for certain type of problems) – offline per day
 - Do not want to manage the big data analytics system
 - Not worry about data privacy/regulation

What would you recommend for solving the requirements?

Example – Legacy then how to deal with big data analytics

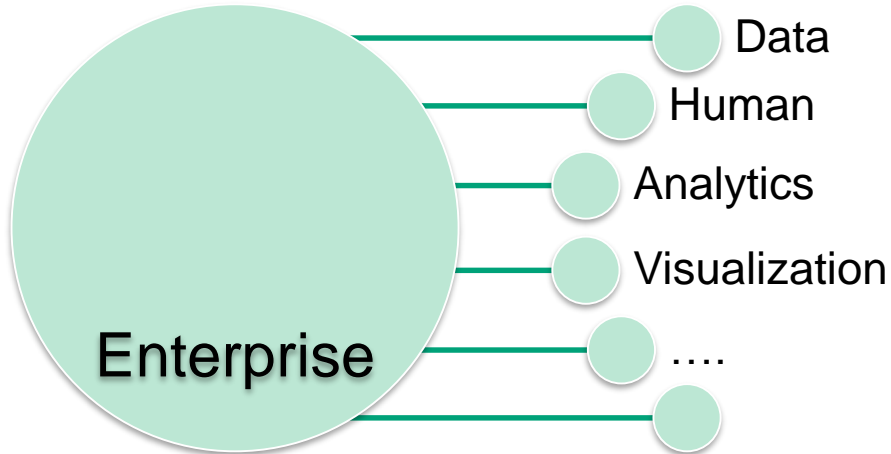


So many types of services from different providers. Anyway to simplify the management of services for the developer/operator?

API MANAGEMENT AND BIG DATA

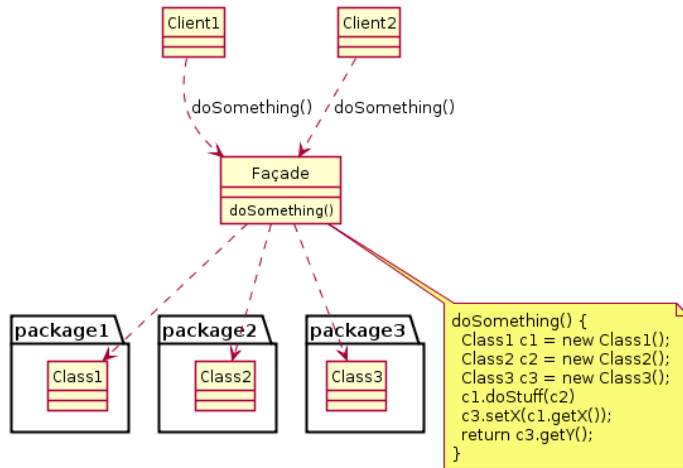
Ecosystem view for advanced service engineering

- Complex data analytics applications → need to understand potential service units from an **ecosystem perspective**
 - Interdependent systems: Social computing, mobile computing, cloud computing, data management, etc.
 - Different functions (analytics, visualization, communications, etc.)
 - Too many different types of customers (and their interactions)
 - Blending vertical and horizontal analytics

- APIs are key! Why?
 - Enable access to data and function from entities in your ecosystem
 - Virtualization
- 
- The diagram illustrates an 'Enterprise' as a large light green circle on the left. Five horizontal lines connect it to a series of smaller light green circles on the right. The circles are labeled from top to bottom: 'Data', 'Human', 'Analytics', 'Visualization', and '....'. A final unlabeled circle is at the bottom, connected by a line.
- An API is an asset
 - We need to have lifecycle, pricing, management, etc.

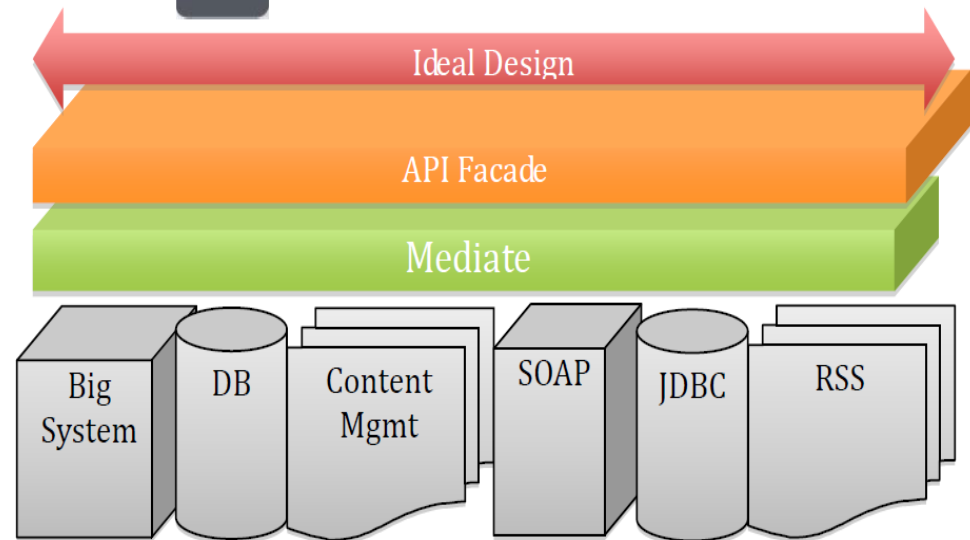
Check <http://www.apiacademy.co> for some useful tutorials

API Fasade

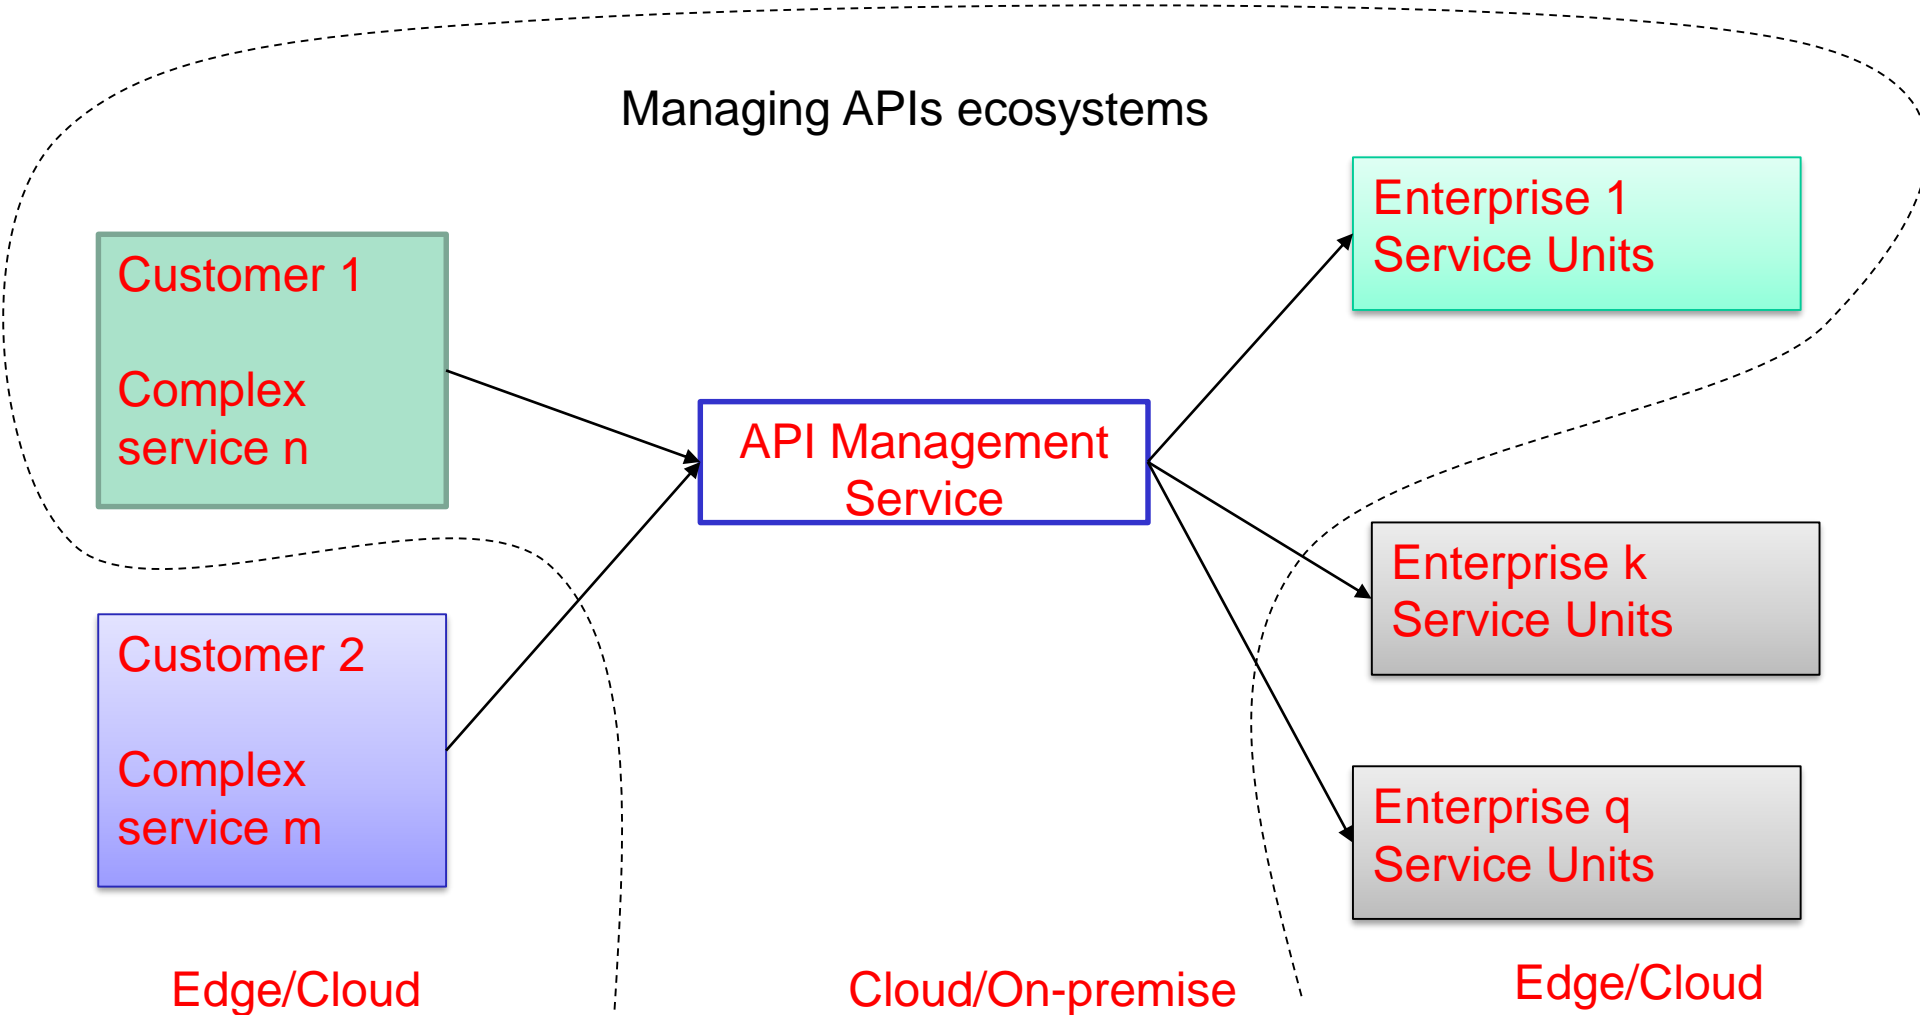


Source:
https://en.wikipedia.org/wiki/Facade_pattern

Source: Web API Design, Brian Mulloy
<http://apigee.com/about/resources/ebooks/web-api-design>



API management & APIs as a service



Development of APIs

- Not just the functions behind the APIs
 - This we have learned since a long time
- Emerging (business/service) management aspects
 - Usage control and security
 - Any where from any device for any customer
 - Interfaces (communications, inputs/output formats)
 - APIs as a service:
 - Availability and reliability of APIs are important – think APIs are similar to a service that your client will consume

Issues on APIs management

- Publish
 - Business and operation planning
 - API usage schemes (e.g., pricing, data concerns)
 - API payload transform policies
 - API throttling
 - API publish and discovery (like service discovery?)
- Management
 - Management roles in enterprises, versions, etc.
- Monitoring and analytics
 - monitoring and analytics information (availability, types of customers, usage frequencies, etc.)
- Microservices and IaC: e.g., with Kubernetes

Some well-known frameworks

- <https://konghq.com/>
- <http://apigee.com>
- Oracle API management:
<http://www.oracle.com/us/products/middleware/soa/api-management/overview/index.html>
- <http://wso2.com/api-management/>
- <http://www.ca.com/us/lpg/layer-7-redirects.aspx>
- <https://www.mashape.com/>
- <http://apiaxle.com/>

Build your own APIs ecosystem

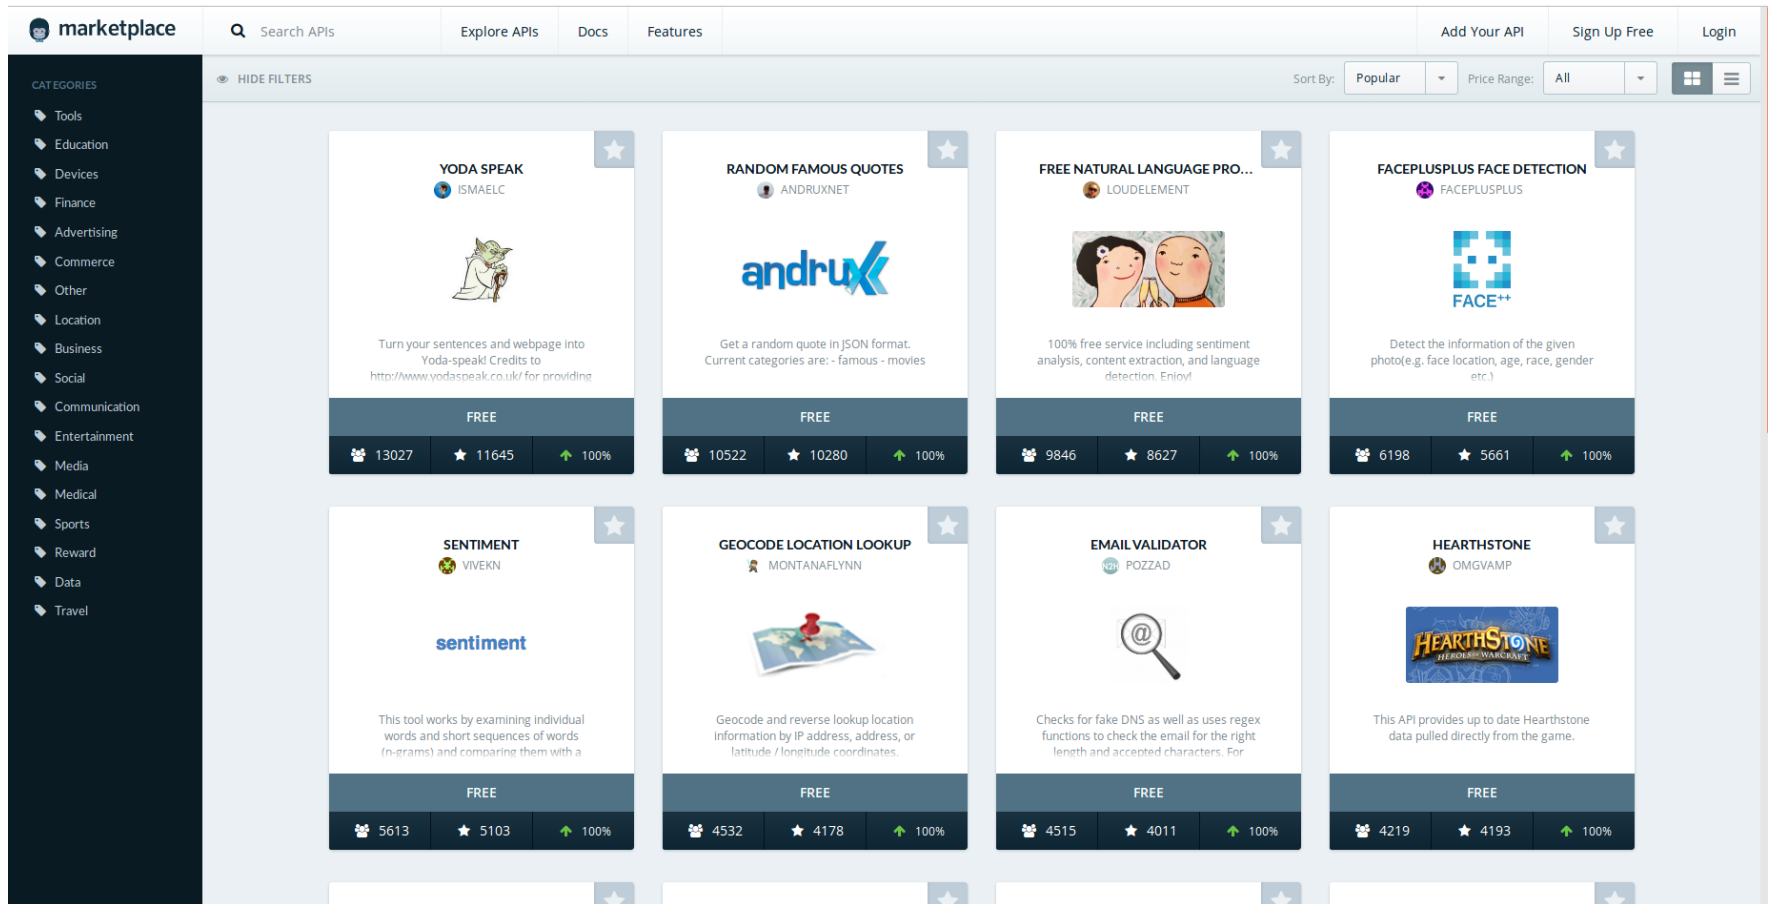
- Which APIs you need? Which ones are crucial for you to build complex services?
 - Data APIs
 - Data collection, Visualization, Analytics APIs
 - Communication
 - Coordination of tasks

→ API management for IoT?

(<http://ubiquity.acm.org/article.cfm?id=2822873>)

- API marketplaces → your APIs
- Using existing API platforms to manage your APIs

Examples of an API marketplace



The screenshot displays a marketplace interface for APIs. The top navigation bar includes a search bar, tabs for 'Explore APIs', 'Docs', and 'Features', and buttons for 'Add Your API', 'Sign Up Free', and 'Login'. A left sidebar lists various categories such as Tools, Education, Devices, Finance, Advertising, Commerce, Other, Location, Business, Social, Communication, Entertainment, Media, Medical, Sports, Reward, Data, and Travel. The main content area shows a grid of API cards, each with a star icon in the top right corner. Each card provides the API name, provider, a brief description, and usage statistics (calls, stars, and success rate).

API Name	Provider	Description	Calls	Stars	Success Rate
YODA SPEAK	ISMAELC	Turn your sentences and webpage into Yoda-speak! Credits to http://www.yodaspeak.co.uk/ for providing	13027	11645	100%
RANDOM FAMOUS QUOTES	ANDRUXNET	Get a random quote in JSON format. Current categories are: - famous - movies	10522	10280	100%
FREE NATURAL LANGUAGE PRO...	LOUDELEMENT	100% free service including sentiment analysis, content extraction, and language detection. Enjoy!	9846	8627	100%
FACEPLUSPLUS FACE DETECTION	FACEPLUSPLUS	Detect the information of the given photo(e.g. face location, age, race, gender etc.)	6198	5661	100%
SENTIMENT	VIVEKN	This tool works by examining individual words and short sequences of words (n-grams) and comparing them with a	5613	5103	100%
GEOCODE LOCATION LOOKUP	MONTANAFLYNN	Geocode and reverse lookup location information by IP address, address, or latitude / longitude coordinates.	4532	4178	100%
EMAIL VALIDATOR	POZZAD	Checks for fake DNS as well as uses regex functions to check the email for the right length and accepted characters. For	4515	4011	100%
HEARTHSTONE	OMGVAMP	This API provides up to date Hearthstone data pulled directly from the game.	4219	4193	100%

Use API Management for your mini project?



WHY API PLATFORM

PRODUCTS

CUSTOMERS

COMMUNITY

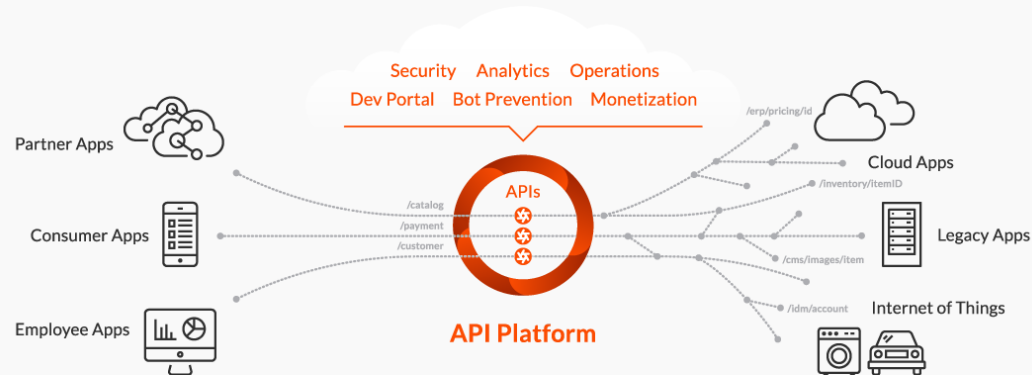
DEVELOPERS

[SIGN UP FREE](#)

[SIGN IN](#)

Manage your APIs anywhere

Design, secure, analyze, and scale your APIs with the Apigee API platform



From <https://apigee.com>

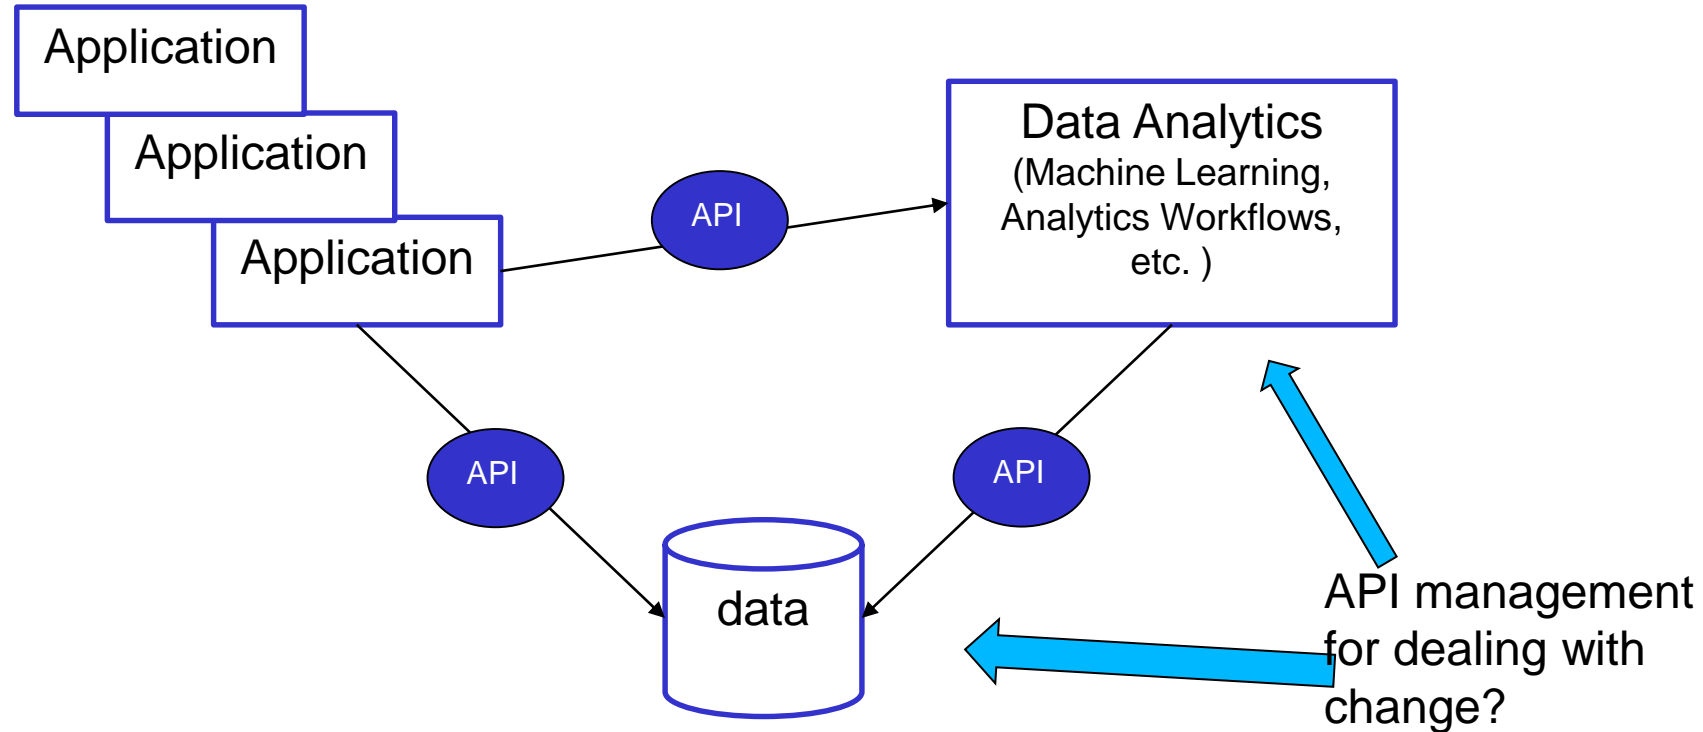
What would be the relationship
between API management and big
data?

Aspects:

- Data access and contract
- New source of data
- Data analytics

Changes in Application, Analytics and data

All are changing internally. Can we keep the API remains and new APIs are added



Example of Architecture Design from Amazon

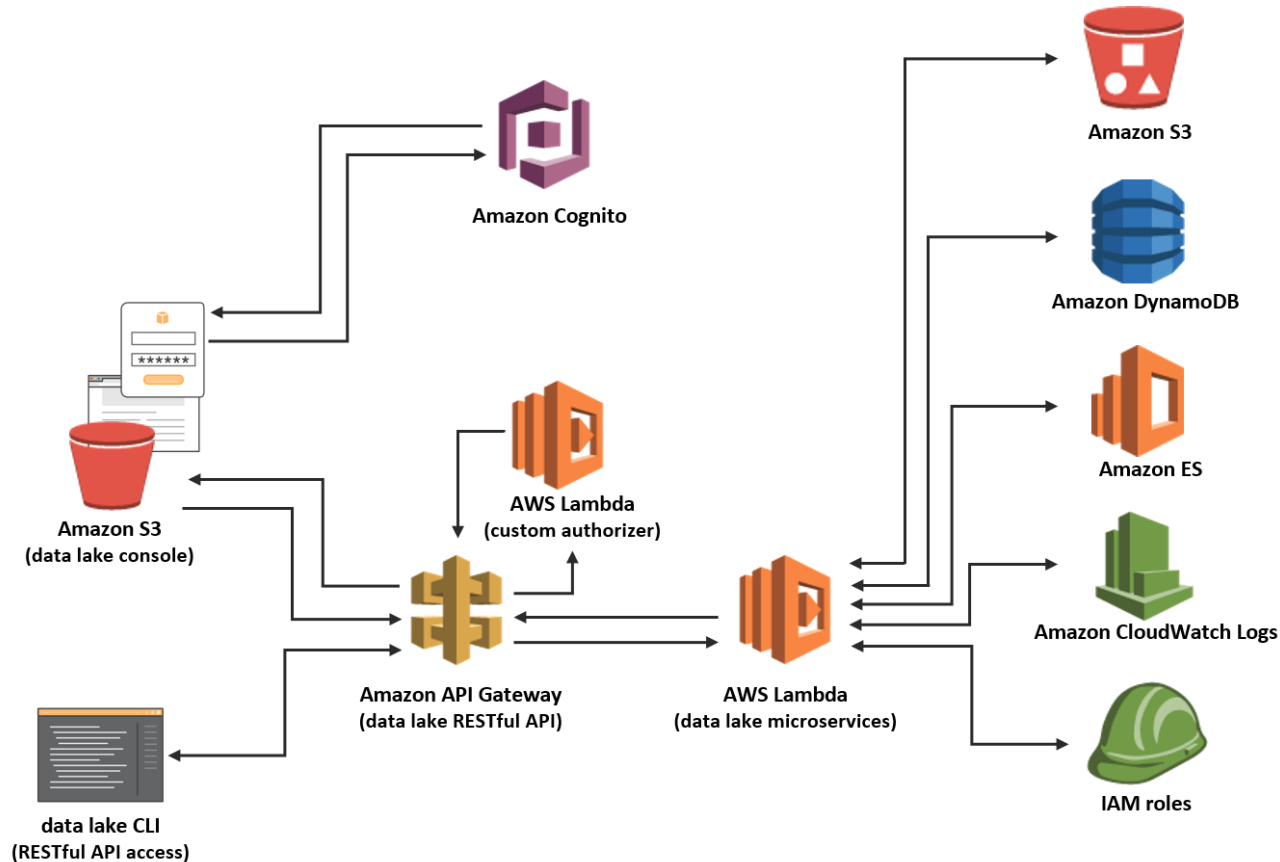


Figure source: <https://aws.amazon.com/answers/big-data/data-lake-solution/>

- Read mentioned papers
- Analyze the relationships between programming models and system infrastructures for data analytics across multiple domains
- Examine <http://cloudcomputingpatterns.org> and see how it supports data analytics patterns
- Mini project:
 - Develop data analytics with suitable frameworks and APIs for your project

Data analytics within a single system

Some papers

1. Andrew Pavlo, Erik Paulson, Alexander Rasin, Daniel J. Abadi, David J. DeWitt, Samuel Madden, and Michael Stonebraker. 2009. A comparison of approaches to large-scale data analysis. In Proceedings of the 2009 ACM SIGMOD International Conference on Management of data (SIGMOD '09), Carsten Binnig and Benoit Dageville (Eds.). ACM, New York, NY, USA, 165-178. DOI=10.1145/1559845.1559865 <http://doi.acm.org/10.1145/1559845.1559865>
2. Leonardo Neumeyer, Bruce Robbins, Anish Nair, Anand Kesari: S4: Distributed Stream Computing Platform. ICDM Workshops 2010: 170-177
3. Jerry Chou, Mark Howison, Brian Austin, Kesheng Wu, Ji Qiang, E. Wes Bethel, Arie Shoshani, Oliver Rübel, Prabhat, and Rob D. Ryne. 2011. Parallel index and query for large scale data analysis. In Proceedings of 2011 International Conference for High Performance Computing, Networking, Storage and Analysis (SC '11). ACM, New York, NY, USA, , Article 30 , 11 pages. DOI=10.1145/2063384.2063424 <http://doi.acm.org/10.1145/2063384.2063424>
4. Boduo Li, Edward Mazur, Yanlei Diao, Andrew McGregor, Prashant J. Shenoy: A platform for scalable one-pass analytics using MapReduce. SIGMOD Conference 2011: 985-996
5. Fabrizio Marozzo, Domenico Talia, Paolo Trunfio: A Cloud Framework for Parameter Sweeping Data Mining Applications. CloudCom 2011: 367-374
6. Yingyi Bu, Bill Howe, Magdalena Balazinska, Michael D. Ernst: HaLoop: Efficient Iterative Data Processing on Large Clusters. PVLDB 3(1): 285-296 (2010)

Thanks for your attention

Hong-Linh Truong
Faculty of Informatics, TU Wien
hong-linh.truong@tuwien.ac.at
www.infosys.tuwien.ac.at/staff/truong
@linhsolar